

Training Pipeline Specification

AI Sales-Enablement Platform

Overview

The training pipeline implements a continuous learning system that processes sales interactions, generates high-quality training data, and continuously improves model performance through feedback loops. The pipeline is designed for scalability, data quality, and rapid iteration cycles.

1. Transcript Ingestion and Preprocessing

Data Sources & Formats

Voice Platform Transcripts:

- **Zoom:** JSON format with speaker identification, timestamps, confidence scores
- **Microsoft Teams:** VTT format with meeting metadata and participant information
- **Gong:** Proprietary format with conversation intelligence annotations
- **Generic:** Support for SRT, WebVTT, and plain text formats

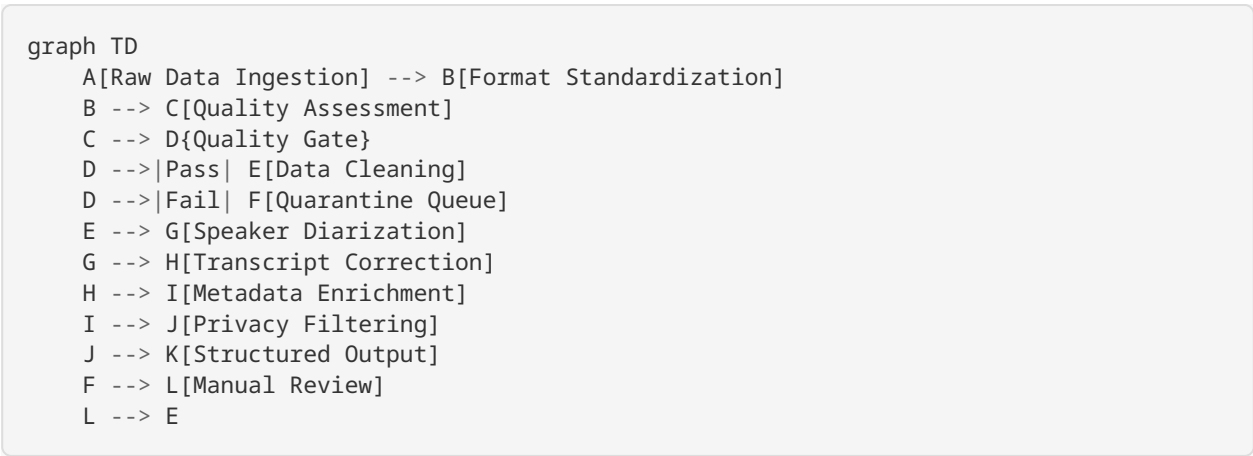
Email Communications:

- **Raw Email:** MIME format with headers, threading information, attachments
- **Email Threads:** Conversation reconstruction with temporal ordering
- **Calendar Context:** Meeting invitations, follow-ups, and scheduling data

CRM Data:

- **Opportunity Records:** Deal stages, amounts, close dates, competitors
- **Contact Information:** Roles, company data, interaction history
- **Activity Logs:** Calls, emails, meetings, and task completions

Preprocessing Pipeline



Data Cleaning & Normalization

Transcript Cleaning:

```
# Example preprocessing steps
def clean_transcript(raw_transcript):
    # Remove filler words and false starts
    cleaned = remove_fillers(raw_transcript, ['um', 'uh', 'like', 'you know'])

    # Correct common transcription errors
    cleaned = apply_domain_corrections(cleaned, sales_terminology_dict)

    # Normalize speaker labels
    cleaned = standardize_speakers(cleaned)

    # Add punctuation and capitalization
    cleaned = apply_punctuation_model(cleaned)

    return cleaned
```

Data Validation:

- **Completeness:** Minimum transcript length, required metadata fields
- **Quality:** Confidence scores, audio quality indicators, speaker clarity
- **Relevance:** Sales conversation detection, business context validation
- **Privacy:** PII detection and redaction, sensitive information filtering

Speaker Diarization & Attribution

Advanced Speaker Identification:

- **Voice Biometrics:** Speaker embedding models for consistent identification
- **Context Clues:** Email signatures, calendar invitations, CRM role mapping
- **Confidence Scoring:** Reliability metrics for speaker attribution
- **Conflict Resolution:** Automated and manual review for ambiguous cases

2. Vector Embedding and Storage Strategy

Embedding Model Architecture

Multi-Modal Embedding Strategy:

- **Text Embeddings:** Sentence-BERT fine-tuned on sales conversations
- **Semantic Embeddings:** Domain-specific models for sales terminology
- **Temporal Embeddings:** Time-aware representations for conversation flow
- **Contextual Embeddings:** CRM data integration for enriched representations

Embedding Generation Pipeline

```
sequenceDiagram
    participant I as Ingestion Service
    participant P as Preprocessing
    participant E as Embedding Service
    participant V as Vector DB
    participant M as Metadata Store

    I->>P: Raw transcript + metadata
    P->>P: Clean and normalize
    P->>E: Processed text chunks
    E->>E: Generate embeddings
    E->>V: Store vectors with IDs
    E->>M: Store metadata + relationships
    V-->>M: Sync vector IDs
```

Chunking Strategy

Hierarchical Chunking:

- **Conversation Level:** Entire call/email thread embeddings
- **Topic Level:** Semantic segmentation based on conversation topics
- **Turn Level:** Individual speaker turns with context windows
- **Sentence Level:** Fine-grained embeddings for precise retrieval

Chunk Optimization:

```
def create_hierarchical_chunks(transcript):
    # Conversation-level chunk
    full_embedding = embed_text(transcript.full_text)

    # Topic-level chunks using semantic segmentation
    topics = segment_by_topic(transcript)
    topic_embeddings = [embed_text(topic.text) for topic in topics]

    # Turn-level chunks with context
    turns = []
    for i, turn in enumerate(transcript.turns):
        context = get_context_window(transcript.turns, i, window_size=3)
        turn_embedding = embed_text(f"{context}\n{turn.text}")
        turns.append(turn_embedding)

    return {
        'full': full_embedding,
        'topics': topic_embeddings,
        'turns': turns
    }
```

Vector Database Design

Database Selection: Pinecone for managed service or Weaviate for self-hosted

Index Configuration:

- **Dimensions:** 768 (sentence-BERT) or 1536 (OpenAI embeddings)
- **Similarity Metric:** Cosine similarity for semantic search
- **Sharding Strategy:** By customer/tenant for data isolation
- **Replication:** Multi-region deployment for availability

Metadata Schema:

```
{
  "id": "unique_chunk_identifier",
  "conversation_id": "parent_conversation_id",
  "chunk_type": "full|topic|turn|sentence",
  "timestamp": "2024-01-15T10:30:00Z",
  "speakers": ["sales_rep", "prospect"],
  "deal_stage": "qualification",
  "deal_value": 50000,
  "industry": "technology",
  "conversation_outcome": "meeting_scheduled",
  "quality_score": 0.85,
  "privacy_level": "internal"
}
```

3. Continuous Learning and Feedback Integration

Feedback Collection System

Multi-Channel Feedback:

- **Direct User Feedback:** Thumbs up/down, corrections, quality ratings
- **Behavioral Signals:** Click-through rates, time spent, actions taken
- **Business Outcomes:** Deal progression, meeting bookings, email responses
- **A/B Testing:** Comparative performance across model versions

Feedback Processing Pipeline

```
graph LR
  A[User Interactions] --> B[Feedback Aggregation]
  B --> C[Quality Validation]
  C --> D[Feedback Scoring]
  D --> E[Training Example Generation]
  E --> F[Incremental Learning]
  F --> G[Model Validation]
  G --> H[Deployment Pipeline]
  H --> I[Performance Monitoring]
  I --> A
```

Training Data Generation

Positive Examples:

- High-rated insights with user confirmations
- Successful conversation patterns leading to positive outcomes
- Effective email templates with high response rates
- Winning sales strategies from closed deals

Negative Examples:

- Low-rated or corrected model outputs
- Conversation patterns from lost deals
- Ineffective communication strategies
- Factually incorrect or irrelevant insights

Data Augmentation:

```
def generate_training_examples(feedback_data):
    positive_examples = []
    negative_examples = []

    for feedback in feedback_data:
        if feedback.rating >= 4: # High quality
            positive_examples.append({
                'input': feedback.query,
                'output': feedback.response,
                'context': feedback.retrieved_docs,
                'weight': feedback.confidence_score
            })
        elif feedback.rating <= 2: # Low quality
            negative_examples.append({
                'input': feedback.query,
                'output': feedback.response,
                'correction': feedback.user_correction,
                'weight': 1.0 - feedback.confidence_score
            })

    return positive_examples, negative_examples
```

Incremental Learning Strategy**LoRA Adapter Updates:**

- **Frequency:** Weekly incremental updates with monthly full retraining
- **Data Requirements:** Minimum 100 high-quality examples per update
- **Validation:** Hold-out test set for performance verification
- **Rollback:** Automatic reversion if performance degrades

Learning Rate Scheduling:

- **Warm-up:** Gradual learning rate increase for stability
- **Decay:** Exponential decay to prevent catastrophic forgetting
- **Adaptive:** Learning rate adjustment based on validation performance

4. Model Evaluation and Drift Detection**Evaluation Framework****Automated Evaluation Metrics:**

```

class ModelEvaluator:
    def __init__(self):
        self.metrics = {
            'relevance': SemanticSimilarityMetric(),
            'accuracy': FactualAccuracyMetric(),
            'coherence': CoherenceMetric(),
            'helpfulness': HelpfulnessMetric()
        }

    def evaluate_model(self, model, test_dataset):
        results = {}
        for metric_name, metric in self.metrics.items():
            score = metric.compute(model, test_dataset)
            results[metric_name] = score

        return results

```

Business Impact Metrics:

- **Sales Performance:** Deal velocity, win rates, average deal size
- **User Adoption:** Feature usage, session duration, return rates
- **Efficiency Gains:** Time saved on manual analysis, faster decision making
- **Quality Improvements:** Reduced errors, increased customer satisfaction

Drift Detection System

Statistical Drift Detection:

```

def detect_embedding_drift(current_embeddings, reference_embeddings):
    # Compute KL divergence between distributions
    kl_div = compute_kl_divergence(current_embeddings, reference_embeddings)

    # Population Stability Index
    psi = compute_psi(current_embeddings, reference_embeddings)

    # Wasserstein distance for distribution comparison
    wasserstein_dist = compute_wasserstein_distance(
        current_embeddings, reference_embeddings
    )

    drift_score = {
        'kl_divergence': kl_div,
        'psi': psi,
        'wasserstein': wasserstein_dist,
        'drift_detected': kl_div > 0.1 or psi > 0.2
    }

    return drift_score

```

Performance Drift Monitoring:

- **Response Quality:** Automated evaluation against golden datasets
- **User Satisfaction:** Trend analysis of feedback scores
- **Business Metrics:** Correlation with sales performance indicators
- **Technical Metrics:** Latency, throughput, error rates

Alerting and Response System

Alert Thresholds:

- **Critical:** >20% performance degradation, immediate model rollback
- **Warning:** 10-20% degradation, accelerated retraining schedule
- **Info:** 5-10% degradation, increased monitoring frequency

Automated Response Actions:

1. **Immediate:** Circuit breaker activation, fallback to previous model
2. **Short-term:** Emergency retraining with recent high-quality data
3. **Medium-term:** Root cause analysis and architecture adjustments
4. **Long-term:** Comprehensive model architecture review

Continuous Improvement Loop

```
graph TB
    A[Performance Monitoring] --> B{Drift Detected?}
    B -->|Yes| C[Root Cause Analysis]
    B -->|No| D[Continue Monitoring]
    C --> E[Data Quality Check]
    E --> F[Retraining Strategy]
    F --> G[Model Update]
    G --> H[A/B Testing]
    H --> I[Performance Validation]
    I --> J{Improvement?}
    J -->|Yes| K[Production Deployment]
    J -->|No| L[Rollback & Investigate]
    K --> A
    L --> C
    D --> A
```

Implementation Timeline

Phase 1: Foundation (Weeks 1-4)

- Data ingestion pipeline setup
- Basic preprocessing and cleaning
- Initial embedding generation
- Vector database deployment

Phase 2: Core Training (Weeks 5-8)

- LoRA fine-tuning pipeline
- Feedback collection system
- Basic evaluation framework
- Initial model deployment

Phase 3: Advanced Features (Weeks 9-12)

- Drift detection system
- Automated retraining
- A/B testing framework
- Performance optimization

Phase 4: Production Hardening (Weeks 13-16)

- Monitoring and alerting
 - Disaster recovery procedures
 - Security hardening
 - Documentation and training
-

Success Metrics

Technical Metrics:

- **Model Performance:** >85% relevance score, <2s inference latency
- **Data Quality:** >90% successful preprocessing, <5% quarantine rate
- **System Reliability:** 99.9% uptime, <1% error rate

Business Metrics:

- **User Adoption:** >80% weekly active users, >4.0 satisfaction score
- **Sales Impact:** 15% improvement in deal velocity, 10% increase in win rate
- **Efficiency:** 30% reduction in manual analysis time

This training pipeline specification provides a comprehensive framework for continuous learning and improvement, ensuring the AI sales-enablement platform evolves with changing business needs and maintains high performance standards.